

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Preprocessing & Model Selection
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import classification_report, confusion_matrix, accuracy_sc

# Classifier
from sklearn.linear_model import LogisticRegression
# You can easily swap this for DecisionTreeClassifier or RandomForestClassifier

# --- Step 1: Load the Data ---
try:
    df = pd.read_csv("D:/Dalmia/Tri-5 2025/ML & DL/ML Internship/Churn_Modelling.csv")
    print("Data loaded successfully.")
except FileNotFoundError:
    print("Error: Make sure 'Churn_Modelling.csv' is in the same directory.")
    # Exit or handle error

print("\nInitial Data Structure:")
print(df.head())
print(df.info())
```

Data loaded successfully.

Initial Data Structure:

```
RowNumber CustomerId Surname CreditScore Geography Gender Age \
0 1 15634602 Hargrave 619 France Female 42
1 2 15647311 Hill 608 Spain Female 41
2 3 15619304 Onio 502 France Female 42
3 4 15701354 Boni 699 France Female 39
4 5 15737888 Mitchell 850 Spain Female 43
```

```
Tenure Balance NumOfProducts HasCrCard IsActiveMember \
0 2 0.00 1 1 1
1 1 83807.86 1 0 1
2 8 159660.80 3 1 0
3 1 0.00 2 0 0
4 2 125510.82 1 1 1
```

EstimatedSalary Exited

```
0 101348.88 1
1 112542.58 0
2 113931.57 1
3 93826.63 0
4 79084.10 0
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10000 entries, 0 to 9999

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	RowNumber	10000	non-null int64
1	CustomerId	10000	non-null int64
2	Surname	10000	non-null object
3	CreditScore	10000	non-null int64
4	Geography	10000	non-null object
5	Gender	10000	non-null object
6	Age	10000	non-null int64
7	Tenure	10000	non-null int64
8	Balance	10000	non-null float64
9	NumOfProducts	10000	non-null int64
10	HasCrCard	10000	non-null int64
11	IsActiveMember	10000	non-null int64
12	EstimatedSalary	10000	non-null float64
13	Exited	10000	non-null int64

dtypes: float64(2), int64(9), object(3)

memory usage: 1.1+ MB

None

```
In [3]: # --- Step 2.1: Drop Irrelevant Columns ---
# RowNumber, CustomerId, and Surname do not contribute to churn prediction
df = df.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1)

# --- Step 2.2: Define Features (X) and Target (y) ---
# The target variable is 'Exited' (1 = Churned, 0 = Not Churned)
X = df.drop('Exited', axis=1)
y = df['Exited']

# --- Step 2.3: Identify Feature Types ---
# 'Geography' and 'Gender' need one-hot encoding.
# All other numerical columns need scaling.
categorical_features = ['Geography', 'Gender']
numerical_features = X.select_dtypes(include=np.number).columns.tolist()
```

```
# Remove features we dropped or the target
for col in ['Exited', 'RowNumber', 'CustomerId', 'Surname']:
    if col in numerical_features:
        numerical_features.remove(col)
```

```
In [5]: # --- Step 3.1: Create Preprocessing Steps ---
# 1. Standard Scaler for numerical columns
numerical_transformer = StandardScaler()

# 2. One-Hot Encoder for categorical columns
categorical_transformer = OneHotEncoder(handle_unknown='ignore')

# Combine transformers using ColumnTransformer
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numerical_transformer, numerical_features),
        ('cat', categorical_transformer, categorical_features)
    ],
    remainder='passthrough' # Keep any other columns if they existed
)

# --- Step 3.2: Create the Full Machine Learning Pipeline ---
# Define the model to use
model = LogisticRegression(random_state=42, solver='liblinear')

# Create the pipeline: Preprocessor -> Model
churn_pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                                 ('classifier', model)])

# --- Step 3.3: Split Data ---
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y # Stratify ensures equal ch
)

print("\nStarting Model Training...")
# Fit the pipeline to the training data
churn_pipeline.fit(X_train, y_train)
print("Training complete!")
```

Starting Model Training...

Training complete!

```
In [7]: # --- Step 4: Make Predictions ---
y_pred = churn_pipeline.predict(X_test)
y_proba = churn_pipeline.predict_proba(X_test)[:, 1]

# --- Step 4.1: Model Performance Metrics ---
print("\n--- Model Evaluation on Test Set ---")

# 1. Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy Score: {accuracy:.4f}")

# 2. ROC-AUC Score (Area Under the Curve - crucial for imbalance)
roc_auc = roc_auc_score(y_test, y_proba)
print(f"ROC-AUC Score: {roc_auc:.4f}")

# 3. Classification Report (Precision, Recall, F1-Score)
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

```
# --- Step 4.2: Confusion Matrix Visualization ---
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted No Churn', 'Predicted Churn'],
            yticklabels=['Actual No Churn', 'Actual Churn'])
plt.title('Confusion Matrix')
plt.ylabel('Actual Label')
plt.xlabel('Predicted Label')
plt.show()
```


--- Model Evaluation on Test Set ---

Accuracy Score: 0.8080

ROC-AUC Score: 0.7748

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.97	0.89	1593
1	0.59	0.19	0.28	407
accuracy			0.81	2000
macro avg	0.71	0.58	0.59	2000
weighted avg	0.78	0.81	0.77	2000



In []: